

Roberto Halpin-Gregorio

rgh224@cornell.edu • (650) 417-3894 • Portfolio: roberto-hg.github.io • LinkedIn: [roberto-hg](#)

Education

Cornell University

M.S. in Computer Science, GPA: 3.98

B.S. in Computer Science, Major GPA: 3.76

Ithaca, NY

Present

May 2020

Experience

Bharath Hariharan Lab

Oct 2020 - Present

Research Assistant — Data Augmentation with Deep Generative Models [PDF](#)

- Built, trained, and tested **deep generative models** to augment image datasets, improving training speed and accuracy by over 8% in multiple machine learning applications
- Augmented vision datasets with more than 10,000 data-efficient **Generative Adversarial Network (GAN)** generated images improving prediction in supervised, semi- and self-supervised learning tasks using **PyTorch**, **TensorFlow**, Pillow, and OpenCV

Madeleine Udell Lab

Apr 2022 - June 2023

Research Assistant — Missing Real-world Data [KDD23 Publication](#)

- Developed novel methods to handle **missing tabular data** in real-world datasets by incorporating informative missing patterns, improving RMSE, AUC, and wall-clock times over standard techniques
- Performed missing data experiments on **real-world data**: MIMIC-III clinical data and several OpenML datasets with linear+logistic regression, XGboost, multilayer perceptron (MLP), and tabular transformer models using PyTorch and **scikit-learn**

Kilian Q. Weinberger Lab

May 2018 - May 2021

Research Assistant — Autonomous Vehicles

- Led the development of an end-to-end **data, training, and inference pipeline** for occluded road segmentation using state-of-the-art **Convolutional Neural Network (CNN)** architectures in PyTorch and the Berkeley Deep Drive (BDD), Cityscapes, and in-house curated datasets
- Constructed a robust Amazon MTurk labeling pipeline with a **custom JavaScript web labeling tool** for road segmentation labeling using **Amazon Web Services (AWS)** cloud storage
- Implemented and tested multiple state-of-the-art **3D trackers and object detectors** in PyTorch and TensorFlow on full sensor datasets – KITTI, NuScenes, Lyft, Waymo, Argo

Cornell Systems Lab

Oct 2019 - May 2020

Research Assistant — Distributed Machine Learning [PDF](#)

- Developed a novel asynchronous **distributed machine learning** scheme for multi-GPU clusters, improving training time by 13% while maintaining statistical performance compared to other distributed methods
- Implemented deep learning models: ResNet, VGG, and random Fourier feature logistic regression with SGD and SVRG in a **distributed GPU parameter server** setting using C++ BLAS, and TensorFlow and PyTorch distributed

John Owens Lab — U.C. Davis

June 2017 - August 2017

Research Assistant — Sparse GPU Optimization

- Designed and implemented **custom NVIDIA GPU kernels** using CUDA and efficient sparse algorithms in C to improve sparse vector-matrix and matrix-matrix multiplication over standard implementations

- Provided real-time assistance to thousands of students during weekly office hours and online course Q&A forums
- Created, tested, and graded assignments/exams; organized and oversaw grading sessions
- Managed and mentored 10+ undergraduate teaching assistants
- Relevant courses taught: Machine Learning, Computer Vision, Principles of Large-Scale Machine Learning, Advanced Machine Learning Systems, Artificial Intelligence

Skills

Languages: Python, C++, Javascript, C

ML Areas: Computer Vision, Deep Learning, Data Augmentation, Training and Inference Pipelining, Distributed Training, Generative Modeling, Self-supervised Learning, Data Cleaning and Preparation, Object Detection, Semantic Segmentation, Image Processing, Tabular Data, Missing and Sparse Data, Tracking

ML Models: CNNs, GANs, MLPs, Transformers, XGBoost, Classic ML/Statistical Models

Libraries/Tools: PyTorch, TensorFlow, NumPy, scikit-learn, Pillow, OpenCV, CUDA, AWS, Amazon MTurk

Other Projects

Visual Self-supervised Learning



Spring 2022

- Designed and implemented a novel masked-image vision transformer and contrastive **self-supervised learning** method that achieves >10% accuracy improvements over previous work on the STL-10 dataset using PyTorch

Representation Learning Theory



Fall 2021

- Discovered new theoretical bounds on the performance of classifiers based on feature representation properties
- Improved generalization of previous bounds in literature and extended bounds to a **k-layer neural network**

Pancreas Tumor Segmentation



Fall 2018

- Designed a novel transfer learning strategy in pancreas tumor segmentation using **CNNs implemented in PyTorch**, improving mean Intersection over Union by 15% on Memorial Sloan Kettering Cancer Center's dataset

Relevant Coursework

ML/AI: Computer Vision, Deep Generative Models, Advanced Machine Learning Systems, Advanced Topics in Machine Learning, Machine Learning Theory, Machine Learning, Artificial Intelligence

Stats: Bayesian Statistics and Data Analysis, Statistical Principles

Other: Mathematics of Data Science, Matrix Computations, Algorithmic Ideas for Speeding Up Optimization